



RCOM

LEIC - 2024/2025

2nd Laboratory Work

Configuration and Study of a Network

Authors:

Bruno Ferreira up202207863@fe.up.pt

Pedro Marinho up202206854@fe.up.pt

Summary

This project involves developing a file download application that utilizes the FTP protocol and TCP/IP for file transfer. It also requires setting up and configuring a computer network to test the application in a real-world environment. The primary goal of the project is to understand how the FTP protocol and TCP/IP work together to enable file transfers over a network. Additionally, it aims to teach how to set up, configure, and troubleshoot a network to support such operations. Throughout the project, we focus on learning how to configure and set up the network to support the application. By the end, we gain practical experience in both developing the application and setting up the network.

Introduction

This project has two main objectives: developing a download application and setting up a network. The download application was created using the FTP (*File Transfer Protocol*) with TCP (*Transmission Control Protocol*) connections through *sockets*. The network setup in the laboratory aims to run the download application using two bridges connected within a *switch*.

The report is divided into three main sections: Part 1 - Development of a download application, Part 2 - Configuration and study of a computer network, and Conclusions. In Part 1, we will discuss the architecture and results of the download application. In Part 2, we will provide a detailed analysis of the network configuration experiments. Finally, we will conclude the report by reflecting on the objectives we achieved.

Part 1 - Development of a download application

Application Architecture

The application's architecture follows a client-server model, implementing the FTP protocol as defined in RFC959 to facilitate the download of a single file from a specified server, supporting URLs with the syntax outlined in RFC1738.

The program's implementation consists of a primary source file, `download.c`, which contains the main logic and entry point, and a header file, `download.h`, used for defining macros and data structures.

1. `download.c`

Usage and Responsibilities:

- Verifies and analyses the given arguments in the command line. Awaits for an URL formatted according to the RFC1738.
- Configures the FTP connection details using `URL` with the info username, password, host, ip, filename and url path.
- Orchestrates the program's instructions in the *main* function, calling the secondary functions to establish the connection, authenticate, guarantee binary format, go into passive mode, transfer the file and close the connection.

2. *download.h*

Data Structures and function implementations

- Data **struct URL**: Stores the information parsed at the beginning of the FTP, preserving the ip, username, password, host, url path and filename.
- Data **struct response**: Stores the message and the code for each response given by the server.
- Includes functions such as `create_socket()`, for establishing the connection, `get_ip()` for decoding the DNS ip, `parse_URL()` to extract the url information, `close_socket()` for closing the connection, `receiveResponse()` to read the servers response, utilizing `readUntilNewLine()` and `readCode()` as helper functions, `showResponse()` to display in the terminal the servers response, `reset_response()` to clean the current responses element, `writeMessage()` to send the command to the server, `calculate_new_port()` to get the new ip and port give by the server, `read_file()` and `getFileSize()` to download the file.

The *Download* program follows a simple and well structured flux, ensuring the system runs smoothly without any issues. The program runs as described below:

1. Initialization:

Parse the given data through the command line, storing the login credentials, the ip, host, filename and the url path.

2. Establish the FTP connection:

The program establishes the connection using the specified port and ip. Once established the connection to the FTP server, the program executes the authentication of the user, sending its credentials.

3. File transfer:

Before the program solicits the server to pass to passive mode, it tells the server to switch to binary mode, getting in the process the file size. After that, having requested the server to be in passive mode, the program receives information from the server to create a new socket for data transfer. Having that socket created, this new connection is utilized to download the requested file, with the specified url.

4. Termination:

After the file transfer conclusion, both created connections are terminated.

Download Report

To run the following program, compile it using *make* or the command `gcc download.c -o download` to generate the binary program and run it by typing `./download <url>` with an URL with this format: `ftp://[<user>:<password>@]<host>/<url-path>`.

We have tested many inputs with anonymous and non anonymous mode, as well as incorrect urls and what not, to ensure the robustness of our program. Dare say the program has shown no signs of inefficiency.

In the following image we show a successful download via our program.

```

pedro@PedrinIPC:/mnt/c/Users/pmiyu/Faculdade/3Ano/1semestre/RCOM/tmp/FEUP-RCOM/proj-2$ ./download ftp://ftp.up.pt/pub/gnu/emacs/elisp-manual-21-2.8.tar.gz
User: anonymous
Password: password
Host: ftp.up.pt
URL Path: pub/gnu/emacs/elisp-manual-21-2.8.tar.gz
Filename: elisp-manual-21-2.8.tar.gz
Host name : mirrors.up.pt
IP Address : 193.137.29.15
Creating socket
Socket created
220-Welcome to the University of Porto's mirror archive (mirrors.up.pt)
220-
220-
220-All connections and transfers are logged. The max number of connections is 200.
220-
220-For more information please visit our website: http://mirrors.up.pt/
220-Questions and comments can be sent to mirrors@upporto.pt
220-
220-
220

Connected to server

user anonymous
331 Please specify the password.

pass password
230 Login successful.

type I
200 Switching to Binary mode.

pasv
227 Entering Passive Mode (193,137,29,15,202,4).

retr pub/gnu/emacs/elisp-manual-21-2.8.tar.gz
150 Opening BINARY mode data connection for pub/gnu/emacs/elisp-manual-21-2.8.tar.gz (2455995 bytes).
Filename: elisp-manual-21-2.8.tar.gz
Receiving file...
[=====] 100.00%
File received and written to elisp-manual-21-2.8.tar.gz
226 Transfer complete.
Transfer completed successfully

quit
221 Goodbye.

```

Part 2 - Configuration and study of a computer network

Exp 1 - Configure an IP Network

The goal of this experiment is to configure the IP addresses of two computers so that they can communicate with each other. We used the `ifconfig` command to configure the IP addresses of both machines. After configuring the `eth0` interfaces and adding the necessary routes, we used the `ping` command to verify the connectivity between them.

-> What are the ARP packets and what are they used for?

ARP (Address Resolution Protocol) packets are used to map a known IP address to a MAC address in a local network. When a device knows the IP address of another device but needs to find its MAC address to send a packet, it sends an ARP request. The device with the matching IP address responds with an ARP reply, providing its MAC address.

-> What are the MAC and IP addresses of ARP packets and why?

In an ARP request packet, the source IP address is the IP address of the requesting device, and the destination IP address is the target IP address that the sender wants to resolve. The source MAC address is the MAC address of the requesting device, and the destination MAC address is initially set to "broadcast" (FF:FF:FF:FF:FF), since the sender does not yet know the MAC address of the target. In an ARP reply, the source and destination roles are reversed: the source MAC address is that of the target device, and the destination MAC address is the original sender.

-> What packets does the ping command generate?

The `ping` command generates ICMP (Internet Control Message Protocol) Echo Request packets to check the availability of a device on the network. If the target device is reachable, it responds with an ICMP Echo Reply packet.

-> What are the MAC and IP addresses of the ping packets?

For a ping request, the source IP address is the IP address of the device sending the ping, and the destination IP address is the IP address of the device being pinged. The source MAC address is the MAC address of the sender's network interface, and the destination MAC address is either the MAC address of the target device (if it is known) or the MAC address of the router (if the target is not directly reachable).

-> How to determine if a receiving Ethernet frame is ARP, IP, ICMP?

To determine the type of a receiving Ethernet frame, you need to check the Ethernet frame's EtherType field:

- If the EtherType is 0x0806, the frame is an ARP packet.
- If the EtherType is 0x0800, the frame is an IP packet.
- If the frame is an IP packet and the protocol field is 0x01, it is an ICMP packet (ICMP is a protocol within IP).

-> How to determine the length of a receiving frame?

To determine the length of a receiving Ethernet frame, check the total frame size. Ethernet frames range from **64 bytes** (minimum) to **1518 bytes** (or **1522 bytes** with VLAN tagging). The **EtherType** field (Ethernet II) indicates the protocol type, while the **Length/Type** field (IEEE 802.3) shows the payload length.

-> What is the loopback interface and why is it important?

The loopback interface is a virtual network interface (often referred to as lo) used by the computer to communicate with itself. The IP address of the loopback interface is typically 127.0.0.1 (or ::1 for IPv6). The loopback interface is important for testing network-related software and services locally without requiring an actual network connection. It allows the system to verify that the network stack is functioning properly.

Exp 2 - Implement two bridges in a switch

The goal of this experiment is to create two virtual LANs (VLANs), bridge110 and bridge111, and assign tux3 and tux4 to the first VLAN (bridge110) and tux2 to the second VLAN (bridge111). We used the ifconfig command to configure the interfaces and the /interface bridge command to first remove the existing ports where the TUX devices were connected. Then, we created two bridges and added the respective ports to each bridge.

-> How to configure bridge110?

To configure a bridge, it was necessary to enter the switch's configuration console and execute the command /interface bridge add name=bridge110. Then, the ports connected to tux3 and tux4 are removed from the default bridge (/interface bridge port remove [find interface=ether1], where 1 is the port identifier) and added to the respective bridges (/interface bridge port add bridge=bridge110 interface=ether1). This process is repeated for bridge111, resulting in two separate subnets.

-> How many broadcast domains are there? How can you conclude it from the logs?

In this experiment, we ran the ping command from tux3 to tux2, and it failed, as expected, since there is no route between the VLANs. Next, we performed a broadcast ping from tux3 using the command ping -b 172.16.110.255, where a response from tux4 was anticipated, as both are in the same subnet.

After this, the same process was repeated from tux2. In this case, no response was received because no device, other than tux2 itself, is configured on bridge111. Thus, it can be concluded that there are two separate broadcast domains with addresses 172.16.110.255 and 172.16.111.255.

Exp 3 - Configure a Router in Linux

The goal of this experiment is to configure tux4 as a router between the two subnets created in the previous experiment, enabling communication between tux3 and tux2. To achieve this, tux4 is connected to Ethernet interface 1, and this interface is added to the subnet where tux2 is located. In this experiment we used the `ifconfig` command to configure the eth2 in the tux4 and then we connected the tux4 to the bridge 111 using the `/interface bridge` command. After that we activated the IP forwarding and deactivated the ICMP echo-ignore-broadcast. Finally we use the `route add -net` command to create the routes in order for the tux2 and tux3 to be connected.

-> **What routes are there in the tuxes? What are their meaning?**

- **tuxY3:**
 - 172.16.110.0/24 (directly connected to eth1 of tuxY3, in the subnet of bridge110).
 - 172.16.111.0/24 (via 172.16.110.254 as the gateway, which is tuxY4's eth1 IP address).
- **tuxY2:**
 - 172.16.111.0/24 (directly connected to eth1 of tuxY2, in the subnet of bridge111).
 - 172.16.110.0/24 (via 172.16.111.253 as the gateway, which is tuxY4's eth2 IP address).
- **tuxY4 (the router):**
 - 172.16.110.0/24 (directly connected to eth1).
 - 172.16.111.0/24 (directly connected to eth2).

The **routes** represent the way each machine knows how to reach another network or device. For instance, tuxY3 knows that to reach devices in 172.16.111.0/24, it must send packets to tuxY4's 172.16.110.254 IP address, and tuxY2 knows that to reach devices in 172.16.110.0/24, it must send packets to tuxY4's 172.16.111.253 IP address.

-> **What information does an entry of the forwarding table contain?**

An entry in the **forwarding table** (also known as the routing table) typically contains:

- **Destination Network:** The target IP address or subnet to which packets are directed (e.g., 172.16.110.0/24).
- **Gateway:** The IP address of the next-hop router that should handle packets destined for that network (e.g., 172.16.110.254 for tuxY3 or 172.16.111.253 for tuxY2).
- **Interface:** The network interface through which the packets should be sent to reach the destination (e.g., eth1 or eth2 for tuxY4).
- **Metric:** The cost of reaching the destination (optional in simple routing configurations).

This table helps the machine know where to forward packets based on their destination IP.

-> **What ARP messages, and associated MAC addresses, are observed and why?**

ARP messages are crucial for resolving IP addresses to MAC addresses, especially when there is communication between different subnets.

- **ARP requests** will be sent when tuxY3 or tuxY2 need to resolve the MAC address of the router (tuxY4) in their subnet:
 - From tuxY3, an ARP request will be sent for 172.16.110.254 (tuxY4's IP on bridge110).
 - From tuxY2, an ARP request will be sent for 172.16.111.253 (tuxY4's IP on bridge111).
- **ARP replies** will be sent by tuxY4 in response, providing its MAC addresses for the respective interfaces:
 - For 172.16.110.254, tuxY4 will reply with the MAC address of eth1.
 - For 172.16.111.253, tuxY4 will reply with the MAC address of eth2.

These ARP messages are observed to ensure that the machines know how to forward packets to the correct network interface.

-> What ICMP packets are observed and why?

During this experiment, **ICMP Echo Request and Echo Reply packets** are observed. These packets are part of the **ping** tests to verify connectivity between tuxY3, tuxY2, and tuxY4:

- **ICMP Echo Request** packets are sent when tuxY3 pings tuxY2 or tuxY4.
- **ICMP Echo Reply** packets are sent back from the destination machine (tuxY2 or tuxY4) if they are reachable.

These ICMP packets are essential for checking if the network is properly configured and if the router (tuxY4) is correctly forwarding packets between the two subnets.

-> What are the IP and MAC addresses associated to ICMP packets and why?

- **ICMP Echo Request:**
 - **Source IP address:** The IP address of the sender (e.g., 172.16.110.1 for tuxY3).
 - **Destination IP address:** The IP address of the target (e.g., 172.16.111.253 for tuxY2, or 172.16.110.254 for tuxY4).
 - **Source MAC address:** The MAC address of the sender's network interface (e.g., tuxY3's eth1 MAC address).
 - **Destination MAC address:** The MAC address of the destination device. If the packet is being routed, this will be the MAC address of the router's relevant interface (e.g., tuxY4's eth1 or eth2).
- **ICMP Echo Reply:**
 - **Source IP address:** The IP address of the machine replying (e.g., 172.16.111.253 for tuxY2, or 172.16.110.254 for tuxY4).
 - **Destination IP address:** The IP address of the machine that sent the original Echo Request.
 - **Source MAC address:** The MAC address of the device sending the reply (e.g., tuxY2's or tuxY4's eth1 or eth2 MAC address).
 - **Destination MAC address:** The MAC address of the machine that initiated the Echo Request (e.g., tuxY3's eth1 MAC address).

Exp 4 - Configure a Commercial Router and Implement NAT

The goal of this experiment is to configure a commercial router and implement Network Address Translation (NAT). This configuration allows multiple devices within a private network to share a single public IP address when accessing external networks, such as the internet. By setting up NAT, the router handles the translation of private IP addresses to a public one and vice versa, ensuring smooth communication between internal devices and external networks. In this experiment, after connecting the router to the switch we used the /interface bridge command to add the router to the bridge111. Then the /ip command was used to configure the router IP address and the necessary routes to all tux devices, so that they are able to communicate with the router.

-> How to configure a static route in a commercial router?

To configure a static route on a commercial router, we connect the S0 serial cable from the Tux4 to the MTik input of the router, allowing us to configure it via GTKTerm. In GTKTerm, we add the route using the command: /ip route add dst-address=172.16.110.0/24 gateway=172.16.111.253.

-> What are the paths followed by the packets, with and without ICMP redirect enabled, in the experiments carried out and why?

- **Without ICMP Redirect Enabled:** When ICMP redirects are disabled, packets are routed according to the static routing table. In this case, packets from tuxY2 going to tuxY3 would go through tuxY4, as specified in the routing table. The path for ICMP packets from tuxY2 to tuxY3 would first hit tuxY4 (the default router for tuxY2), which will forward the packets to tuxY3 based on its own routing table.
- **With ICMP Redirect Enabled:** ICMP redirects inform the sender that there is a better route available to reach the destination directly. When ICMP redirects are enabled, tuxY2 may receive a redirect message from tuxY4 or RC, instructing it to send packets directly to tuxY3, instead of going through tuxY4. This reduces unnecessary hops and optimizes routing for subsequent packets.

-> How to configure NAT in a commercial router?

To configure NAT (Network Address Translation) on a commercial router, follow these general steps:

- **Access the Router Configuration Interface:** Login via the CLI or web interface.
- **Identify the Interfaces:** Determine which interface is connected to the internal network (private IP range) and which one connects to the external network (public IP range).
- **Enable NAT:**
 - For source NAT (SNAT), configure the router to translate the internal (private) IP addresses to the router's public IP address when accessing external networks.
 - For destination NAT (DNAT), configure the router to translate external IP addresses back to internal private IP addresses for incoming connections.

This would enable NAT for outgoing traffic from the internal network to the external network, we use the : /ip firewall nat enable command.

-> What does NAT do?

NAT (Network Address Translation) allows multiple devices on a private network to share a single public IP address for internet access. It modifies IP addresses in packet

headers, translating private IPs to the router's public IP when leaving the network and vice versa when returning. NAT also enhances security by hiding internal IP addresses from external networks, preventing direct access to internal devices.

-> What happens when tuxY3 pings the FTP server with the NAT disabled? Why?

When NAT is disabled, tuxY3 will attempt to access the FTP server using its private IP address. Since the NAT is no longer translating the private IP address into the router's public IP address, the packets destined for the FTP server might not be routed correctly, especially if the FTP server is outside the private network.

Without NAT, tuxY3's private IP address may not be recognized or routable by the external network, and thus tuxY3 will not be able to communicate with the FTP server. The communication would fail, as the network's routing tables would not know how to handle the private IP address outside the local network.

Disabling NAT essentially breaks the ability of devices in the private network to access external resources because there is no translation between private and public IP addresses.

Exp 5 - DNS

The goal of this experiment is to configure the DNS (Domain Name System) on tux3, tux4, and tux2. A DNS server is used to translate hostnames into their corresponding IP addresses. In this case, the server services.netlab.fe.up.pt contains a database with public IP addresses and their associated hostnames. This configuration enables the systems to resolve domain names to IP addresses, allowing devices to communicate using human-readable names instead of numeric IP addresses.

-> How to configure the DNS service in a host?

To configure the DNS service, it is necessary to edit the resolv.conf file located at /etc/resolv.conf on the tux host. This file needs to be edited to include the following information: echo "nameserver 10.227.20.3" > /etc/resolv.conf. After this configuration, it will be possible to access the internet from the tux machines.

-> What packets are exchanged by DNS and what information is transported?

DNS exchanges two types of packets: the **DNS query packet** and the **DNS response packet**. The **DNS query packet** is sent from the client to the DNS server and contains the hostname (e.g., "www.google.com") that the client wants to resolve into an IP address. The **DNS response packet** is sent from the server to the client and contains the resolved IP address for the queried hostname. This exchange allows the client to convert a human-readable hostname into a machine-readable IP address.

Exp 6 - TCP connections

Once the network was fully configured, the download application developed in Part 1 was used as a case study. The test was successfully conducted by transferring a file from an FTP server, and the application performed as expected.

-> How many TCP connections are opened by your FTP application?

The FTP application typically opens **two TCP connections**:

- **Control connection** (usually on port 21) for sending commands and receiving responses.

- **Data connection** (on a dynamically assigned port) for transferring the actual file data.

-> In what connection is transported the FTP control information?

The **FTP control information** is transported over the **control connection** (typically on TCP port 21). This connection is used to send FTP commands such as USER, PASS, TYPE, PASV, RETR (retrieve) and QUIT as well as to send responses from the server.

-> What are the phases of a TCP connection?

1. **Connection Establishment** (Three-way handshake):
 - The client sends a **SYN** packet to initiate the connection.
 - The server responds with a **SYN-ACK** packet.
 - The client sends a **final ACK** to confirm the connection.
2. **Data Transfer**:
 - After the connection is established, data is transferred between the client and server using the reliable TCP stream.
3. **Connection Termination** (Four-way handshake):
 - The client or server sends a **FIN** packet to terminate the connection.
 - The other party responds with an **ACK** to acknowledge receipt of the FIN.
 - The closing party then sends a **FIN** to fully close the connection.
 - The final **ACK** packet is exchanged to confirm termination.

-> How does the ARQ TCP mechanism work? What are the relevant TCP fields? What relevant information can be observed in the logs?

The **ARQ (Automatic Repeat reQuest)** mechanism in TCP ensures reliable data transmission by using **acknowledgments** and **retransmissions**.

- Relevant TCP fields include:
 - **Sequence Number**: Identifies the position of the first byte of data in the segment.
 - **Acknowledgment Number**: Indicates the next byte expected from the other side.
 - **Flags**: Especially **ACK**, indicating acknowledgment.
 - **Window Size**: The amount of data the receiver is willing to accept.

In the logs, you can observe retransmitted packets, acknowledgment packets, and the flow of data. If there is a loss, the ARQ mechanism will trigger retransmission.

-> How does the TCP congestion control mechanism work? What are the relevant fields. How did the throughput of the data connection evolve along the time? Is it according to the TCP congestion control mechanism?

The **TCP congestion control mechanism** works by dynamically adjusting the data transfer rate to avoid overloading the network. It uses algorithms such as **Slow Start**, **Congestion Avoidance**, **Fast Retransmit**, and **Fast Recovery** to adjust the window size and retransmit lost packets.

- Relevant TCP fields:
 - **Window Size**: This is adjusted according to the network's congestion status.

- **Congestion Window (cwnd):** Determines how much data can be sent before receiving an acknowledgment.

Over time, the throughput typically starts small (slow start), increases gradually (congestion avoidance), and may decrease if packet loss occurs (indicating congestion). In Wireshark, you can observe the window size and retransmissions, which are indicators of the congestion control mechanism at work.

The throughput should follow the TCP congestion control curve, starting slow and increasing over time until it reaches a point where packet loss may reduce the rate.

-> Is the throughput of a TCP data connections disturbed by the appearance of a second TCP connection? How?

Yes, the throughput of a TCP data connection can be disturbed by the appearance of a second TCP connection, particularly if both connections share the same network resources (bandwidth). When the second connection starts, the available bandwidth is divided between both, which can lead to a **reduction in throughput** for the first connection. This happens because TCP congestion control mechanisms try to avoid network congestion by adjusting the window size based on the available bandwidth, causing the throughput to decrease as more traffic is introduced into the network. This effect is more noticeable if both connections are competing for the same network path.

Conclusions

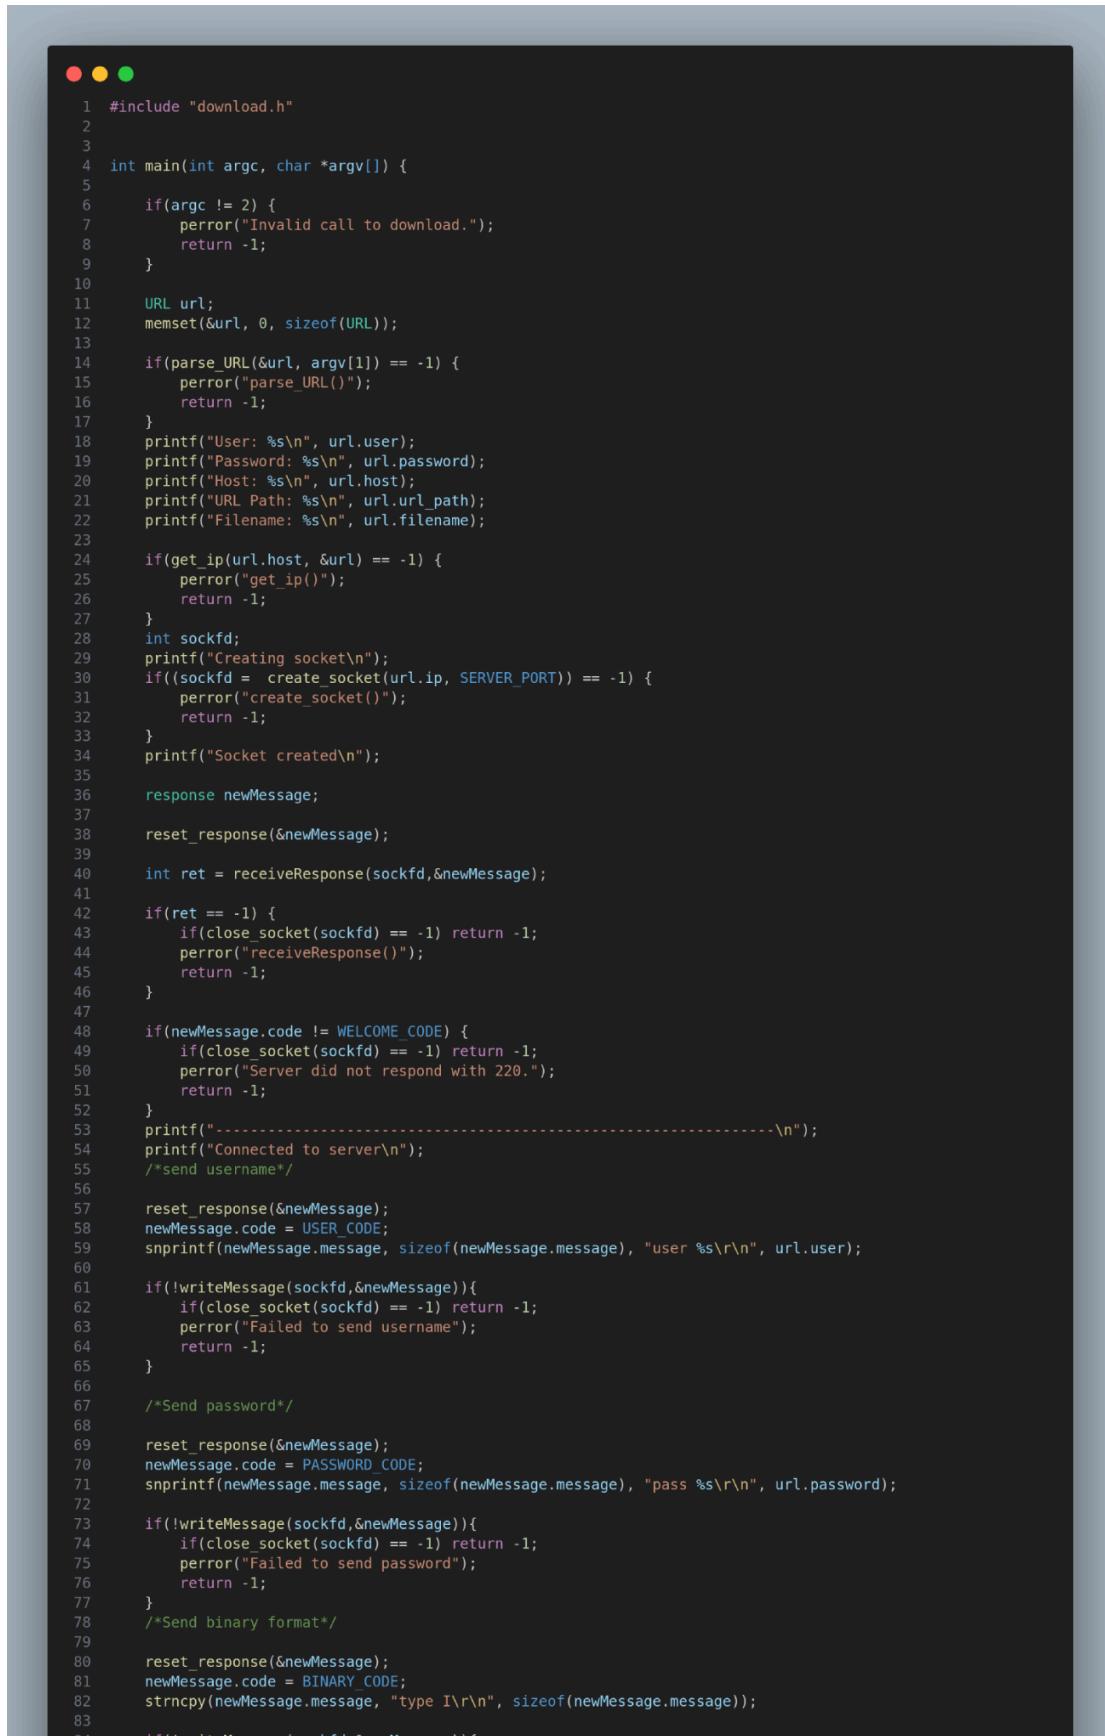
In conclusion, we successfully met the objectives of this project by developing a file download application using FTP and TCP/IP, and setting up a network to test it. Through the development process, we gained valuable insights into how FTP and TCP/IP protocols enable file transfers. Configuring and troubleshooting the network further enhanced our understanding of networking concepts. Overall, the project allowed us to strengthen our skills in both application development and network setup, providing practical experience and laying a solid foundation for future work in software development and networking.

Appendix I - Source Code

download.h

```
● ● ●
1 #ifndef download_H
2 #define download_H
3
4
5 #include <stdio.h>
6 #include <stdlib.h>
7 #include <netdb.h>
8 #include <netinet/in.h>
9 #include <arpa/inet.h>
10 #include <sys/socket.h>
11 #include <unistd.h>
12 #include <string.h>
13 #include <regex.h>
14 #include <fcntl.h>
15
16 #define MAX_LEN 500
17
18 #define SERVER_PORT 21
19 #define CODE_SIZE 4
20 #define WELCOME_CODE 220
21 #define USER_CODE 331
22 #define PASSWORD_CODE 230
23 #define PASSIVE_CODE 227
24 #define MIN_RETREIVE_CODE 100
25 #define MAX_RETREIVE_CODE 199
26 #define RETREIVE_CODE_WITH_SIZE 150
27 #define MIN_TRANSFER_CODE 200
28 #define BINARY_CODE 200
29 #define GOODBYE_CODE 221
30
31 typedef struct
32 {
33     char ip[16];
34     char user[MAX_LEN];
35     char password[MAX_LEN];
36     char host[MAX_LEN];
37     char url_path[MAX_LEN];
38     char filename[MAX_LEN];
39 } URL;
40
41 typedef struct
42 {
43     char message[1024];
44     int code;
45 } response;
46
47
48
49 int create_socket(char *ip,int port);
50 int get_ip(char* hostname, URL *url);
51 int parse_URL(URL *url,char *url_str);
52 int close_socket(int sockfd);
53 int receiveResponse(int socketfd, response *res);
54 int readUntilNewline(int socketfd, char *buf);
55 int readCode(int socketfd, char *code);
56 void showResponse(response *res);
57 void reset_response(response *newMessage);
58 int writeMessage(int sockfd, response *message);
59 int calculate_new_port(char *passiveMsg, URL url);
60 int readfile(int sockfd, char *filename, long long file_size);
61 long long getFileSize(char * message);
62
63 #endif
```

download.c



```
1 #include "download.h"
2
3
4 int main(int argc, char *argv[]) {
5
6     if(argc != 2) {
7         perror("Invalid call to download.");
8         return -1;
9     }
10
11     URL url;
12     memset(&url, 0, sizeof(URL));
13
14     if(parse_URL(&url, argv[1]) == -1) {
15         perror("parse_URL()");
16         return -1;
17     }
18     printf("User: %s\n", url.user);
19     printf("Password: %s\n", url.password);
20     printf("Host: %s\n", url.host);
21     printf("URL Path: %s\n", url.url_path);
22     printf("Filename: %s\n", url.filename);
23
24     if(get_ip(url.host, &url) == -1) {
25         perror("get_ip()");
26         return -1;
27     }
28     int sockfd;
29     printf("Creating socket\n");
30     if((sockfd = create_socket(url.ip, SERVER_PORT)) == -1) {
31         perror("create_socket()");
32         return -1;
33     }
34     printf("Socket created\n");
35
36     response newMessage;
37
38     reset_response(&newMessage);
39
40     int ret = receiveResponse(sockfd,&newMessage);
41
42     if(ret == -1) {
43         if(close_socket(sockfd) == -1) return -1;
44         perror("receiveResponse()");
45         return -1;
46     }
47
48     if(newMessage.code != WELCOME_CODE) {
49         if(close_socket(sockfd) == -1) return -1;
50         perror("Server did not respond with 220.");
51         return -1;
52     }
53     printf("-----\n");
54     printf("Connected to server\n");
55     /*send username*/
56
57     reset_response(&newMessage);
58     newMessage.code = USER_CODE;
59     sprintf(newMessage.message, sizeof(newMessage.message), "user %s\r\n", url.user);
60
61     if(!writeMessage(sockfd,&newMessage)){
62         if(close_socket(sockfd) == -1) return -1;
63         perror("Failed to send username");
64         return -1;
65     }
66
67     /*Send password*/
68
69     reset_response(&newMessage);
70     newMessage.code = PASSWORD_CODE;
71     sprintf(newMessage.message, sizeof(newMessage.message), "pass %s\r\n", url.password);
72
73     if(!writeMessage(sockfd,&newMessage)){
74         if(close_socket(sockfd) == -1) return -1;
75         perror("Failed to send password");
76         return -1;
77     }
78     /*Send binary format*/
79
80     reset_response(&newMessage);
81     newMessage.code = BINARY_CODE;
82     strncpy(newMessage.message, "type I\r\n", sizeof(newMessage.message));
83
84     if(!writeMessage(sockfd,&newMessage)){
```

```

84     if(!writeMessage(sockfd,&newMessage)){
85         if(close_socket(sockfd) == -1) return -1;
86         perror("Failed to pass binary format");
87         return -1;
88     }
89
90     /*Send passive mode*/
91
92     reset_response(&newMessage);
93     newMessage.code = PASSIVE_CODE;
94     strncpy(newMessage.message, "pasv\r\n", sizeof(newMessage.message));
95
96     if(!writeMessage(sockfd,&newMessage)){
97         if(close_socket(sockfd) == -1) return -1;
98         perror("Failed to activate passivemode");
99         return -1;
100    }
101
102    int port2;
103
104    if((port2 = calculate_new_port(newMessage.message,url)) == -1){
105        if(close_socket(sockfd) == -1) return -1;
106        perror("Failed to calculate port");
107        return -1;
108    }
109    int sockfd2;
110    if ((sockfd2 = create_socket(url.ip,port2)) == -1){
111        perror("Failed to create socket");
112        return -1;
113    }
114
115    /*Send retrive*/
116    reset_response(&newMessage);
117    newMessage.code = RETREIVE_CODE_WITH_SIZE;
118    sprintf(newMessage.message, sizeof(newMessage.message), "retr %s\r\n", url.url_path);
119
120    if(!writeMessage(sockfd,&newMessage)){
121        if(close_socket(sockfd) == -1) return -1;
122        perror("Failed to retrive");
123        return -1;
124    }
125    long long fileSize = 1;
126    printf("Filename: %s\n",url.filename);
127    if (newMessage.code == RETREIVE_CODE_WITH_SIZE) {
128        fileSize = getFileSize(newMessage.message);
129    }
130
131    if(fileSize == -1){
132        perror("Failed to get the file size");
133        return -1;
134    }
135
136    if(readfile(sockfd2,url.filename,fileSize) == -1){
137        if(close_socket(sockfd) == -1) return -1;
138        return -1;
139    }
140
141    reset_response(&newMessage);
142
143    ret = receiveResponse(sockfd,&newMessage);
144
145    if(ret == -1) {
146        if(close_socket(sockfd) == -1) return -1;
147        perror("receiveResponse()");
148        return -1;
149    }
150
151    if(newMessage.code < MIN_TRANSFER_CODE) {
152        if(close_socket(sockfd) == -1) return -1;
153        perror("Transfer went wrong.");
154        return -1;
155    }
156
157    printf("Transfer completed successfully\n");
158
159    reset_response(&newMessage);
160    newMessage.code = GOODBYE_CODE;
161    strncpy(newMessage.message, "quit\r\n", sizeof(newMessage.message));
162
163    if(!writeMessage(sockfd,&newMessage)){
164        if(close_socket(sockfd) == -1) return -1;
165        perror("Failed to say goodbye");
166        return -1;
167    }
168
169    if(close_socket(sockfd) == -1) return -1;
170    if(close_socket(sockfd2) == -1) return -1;
171    return 0;

```

```

172 }
173
174 void reset_response(response *res) {
175     if (res == NULL) return; // Safety check
176     res->code = 0; // Reset the response code
177     memset(res->message, 0, sizeof(res->message)); // Clear the message buffer
178 }
179
180 int create_socket(char *ip,int port) {
181
182     int sockfd;
183     struct sockaddr_in server_addr;
184
185     /*server address handling*/
186     bzero((char *) &server_addr, sizeof(server_addr));
187     server_addr.sin_family = AF_INET;
188     server_addr.sin_addr.s_addr = inet_addr(ip); /*32 bit Internet address network byte ordered*/
189     server_addr.sin_port = htons(port); /*server TCP port must be network byte ordered */
190
191     /*open a TCP socket*/
192     if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
193         perror("socket()");
194         return -1;
195     }
196     /*connect to the server*/
197     if (connect(sockfd,
198                 (struct sockaddr *) &server_addr,
199                 sizeof(server_addr)) < 0) {
200         perror("connect()");
201         return -1;
202     }
203
204
205
206     return sockfd;
207 }
208
209
210 int get_ip(char* hostname, URL *url) {
211     struct hostent *h;
212
213     if ((h = gethostbyname(hostname)) == NULL) {
214         perror("gethostbyname()");
215         return -1;
216     }
217
218     printf("Host name : %s\n", h->h_name);
219     printf("IP Address : %s\n", inet_ntoa(*((struct in_addr *) h->h_addr)));
220     strcpy(url->ip, inet_ntoa(*((struct in_addr *) h->h_addr)), sizeof(url->ip) - 1);
221     return 0;
222 }
223
224
225 int parse_URL(URL *url, char *url_str) {
226     // Regular expression to match the URL format
227     const char *regex_pattern = "^(ftp://|(([:@/]+)(:[^@/]+)?@)?([^.]+)(/.+)$";
228
229     regex_t regex;
230     regmatch_t matches[7]; // There are 7 capture groups in the regex
231
232     // Compile the regular expression
233     if (regcomp(&regex, regex_pattern, REG_EXTENDED) != 0) {
234         fprintf(stderr, "Failed to compile regex.\n");
235         return -1;
236     }
237
238     // Execute the regular expression on the URL string
239     if (regexec(&regex, url_str, 7, matches, 0) != 0) {
240         fprintf(stderr, "Invalid URL format.\n");
241         regfree(&regex);
242         return -1;
243     }
244
245     // Extract user and password (if present)
246     if (matches[1].rm_so != -1) {
247         // If user and password are present
248         if (matches[2].rm_so != -1) {
249             strcpy(url->user, url_str + matches[2].rm_so, matches[2].rm_eo - matches[2].rm_so);
250             url->user[matches[2].rm_eo - matches[2].rm_so] = '\0';
251         } else {
252             strcpy(url->user, "anonymous", sizeof(url->user) - 1);
253         }
254
255         if (matches[4].rm_so != -1) {
256             strcpy(url->password, url_str + matches[4].rm_so, matches[4].rm_eo - matches[4].rm_so);
257             url->password[matches[4].rm_eo - matches[4].rm_so] = '\0';
258         } else {
259             strcpy(url->password, "password", sizeof(url->password) - 1); // No password

```

```

259         strncpy(url->password, "password", sizeof(url->password) - 1); // No password
260     }
261 } else {
262     strncpy(url->user, "anonymous", sizeof(url->user) - 1); // No user
263     strncpy(url->password, "password", sizeof(url->password) - 1); // No password
264 }
265
266 // Extract host
267 if (matches[5].rm_so != -1) {
268     strncpy(url->host, url_str + matches[5].rm_so, matches[5].rm_eo - matches[5].rm_so);
269     url->host[matches[5].rm_eo - matches[5].rm_so] = '\0';
270 } else {
271     url->host[0] = '\0'; // Handle missing host
272 }
273
274 // Extract URL path
275 if (matches[6].rm_so != -1) {
276     strncpy(url->url_path, url_str + matches[6].rm_so, matches[6].rm_eo - matches[6].rm_so);
277     url->url_path[matches[6].rm_eo - matches[6].rm_so] = '\0';
278 } else {
279     url->url_path[0] = '\0'; // Handle missing URL path
280 }
281
282 // Extract filename from the path (if exists)
283 char *last_slash = strrchr(url->url_path, '/');
284 if (last_slash && *(last_slash + 1) != '\0') {
285     // There is a filename after the last slash
286     strncpy(url->filename, last_slash + 1, MAX_LEN - 1);
287     url->filename[MAX_LEN - 1] = '\0';
288 } else {
289     // No filename, use the entire path as the filename
290     strncpy(url->filename, url->url_path, MAX_LEN - 1);
291     url->filename[MAX_LEN - 1] = '\0';
292 }
293
294 // Set IP to an empty string (no parsing logic for IP)
295 url->ip[0] = '\0';
296
297 // Free the regex memory
298 regfree(&regex);
299
300     return 0;
301 }
302
303 int close_socket(int sockfd) {
304     if(close(sockfd) < 0) {
305         perror("Error closing socket");
306         return -1;
307     }
308     return 0;
309 }
310
311
312 int receiveResponse(int socketfd, response *res) {
313     int is_num = 1;
314     while (1) {
315
316         if (readUntilNewline(socketfd, res->message) == -1) return -1; // Error reading until newline
317
318         showResponse(res);
319
320         // check if its the code is a number
321         for (int i = 0; i < 3; i++) {
322             if (!(res->message[i] >= '0' && res->message[i] <= '9')) {
323                 is_num = 0;
324                 break;
325             }
326         }
327
328         // Check if its only the code without -
329         if ((res->message[3] == ' ') && is_num) {
330             res->code = atoi(res->message); // Convert the code to an integer
331             return 0;
332         }
333
334         reset_response(res);
335         // Read until a newline character is encountered
336         is_num = 1;
337     }
338
339     return 0;
340 }
341
342 int readCode(int socketfd, char *code) {
343     for (int j = 0; j < CODE_SIZE; j++) {
344         int res = read(socketfd, &code[j], 1);
345         if (res < 0) return -1; // Handle read error
346     }

```

```

347     return 0; // Successfully read the code
348 }
349
350 // Function to read data until a newline character is encountered
351 int readUntilNewline(int sockfd, char *buf) {
352     int res;
353     int i = 0;
354     do {
355         res = read(sockfd, &buf[i++], 1);
356         if (res < 0) return -1; // Handle read error
357     } while (buf[i - 1] != '\n');
358
359     return 0; // Successfully read until newline
360 }
361
362
363 // Message formating is weird
364 void showResponse(response *res) {
365     printf("%s", res->message);
366 }
367
368
369 int writeMessage(int sockfd, response *message){
370     if (sockfd < 0) {
371         perror("Invalid socket descriptor");
372         return -1;
373     }
374
375     if (strlen(message->message) == 0) {
376         fprintf(stderr, "Invalid message to write\n");
377         return -1;
378     }
379
380     printf("\n%s", message->message);
381
382     size_t bytes_sent = 0;
383     size_t to_send = strlen(message->message);
384
385     // Write the message in chunks if necessary
386     while (bytes_sent < to_send) {
387         ssize_t result = write(sockfd, message->message + bytes_sent, to_send - bytes_sent);
388         if (result < 0) {
389             perror("Error writing to socket");
390             return -1;
391         }
392         bytes_sent += result;
393     }
394
395
396     // Receive the server's response
397     response res;
398     memset(&res, 0, sizeof(res));
399
400     if (receiveResponse(sockfd, &res) == -1) {
401         perror("receiveResponse failed");
402         return -1;
403     }
404
405
406     strncpy(message->message, res.message, sizeof(message->message) - 1);
407     message->message[sizeof(message->message) - 1] = '\0';
408
409     if (message->code == RETREIVE_CODE_WITH_SIZE) {
410         message->code = res.code;
411         return res.code >= MIN_RETREIVE_CODE && res.code <= MAX_RETREIVE_CODE;
412     }
413
414     return res.code == message->code;
415 }
416
417 int calculate_new_port(char *passiveMsg, URL url){
418     int ip1, ip2, ip3, ip4, port1, port2;
419
420     int parsed = sscanf(passiveMsg, "%*[^\n](%d,%d,%d,%d,%d)", &ip1, &ip2, &ip3, &ip4, &port1, &port2);
421
422     if (parsed != 6) {
423         perror("Failed to parse passive message");
424         return -1;
425     }
426
427     // Compare the extracted IP with the provided IP (url.ip should be in the same format as the parsed IP)
428     // Assuming url.ip is in the format "ip1.ip2.ip3.ip4"
429     int url_ip1, url_ip2, url_ip3, url_ip4;
430     if (sscanf(url.ip, "%d.%d.%d.%d", &url_ip1, &url_ip2, &url_ip3, &url_ip4) != 4) {
431         perror("Failed to parse URL IP");
432         return -1;
433     }
434 }
```

```

435 // If the IPs don't match, return an error
436 if (ip1 != url_ip1 || ip2 != url_ip2 || ip3 != url_ip3 || ip4 != url_ip4) {
437     perror("IP addresses do not match");
438     return -1;
439 }
440
441 int port = port1*256 + port2;
442
443 return port;
444 }
445
446
447 int readfile(int sockfd, char *filename, long long file_size) {
448     int file_fd;
449     off_t bytes_received = 0;
450     ssize_t res;
451
452     // Open the file in write-only mode. Create it if it doesn't exist.
453     // Set permissions to rw-r--r-- (user can read/write, group and others can read).
454     if ((file_fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, 0644)) < 0) {
455         perror("Open file");
456         return -1;
457     }
458
459     // Buffer for receiving data from the socket
460     char buf[1024]; // Larger buffer to read multiple bytes at once
461
462     printf("Receiving file...\n");
463
464     // Loop to read from the socket and write to the file
465     while (1) {
466         // Read data from the socket
467         res = read(sockfd, buf, sizeof(buf));
468
469         if (res == 0) {
470             // End of transmission
471             break;
472         }
473         if (res < 0) {
474             // Error reading from the socket
475             perror("Failed to read from socket");
476             close(file_fd);
477             return -1;
478         }
479
480         // Write the received data into the file
481         if (write(file_fd, buf, res) < 0) {
482             // Error writing to the file
483             perror("Failed to write to file");
484             close(file_fd);
485             return -1;
486         }
487
488         // Update bytes received
489         bytes_received += res;
490         // Print loading bar based on the total file size
491         float progress = (float)bytes_received / file_size * 100; // Progress based on actual file size
492         int bar_width = 50; // Width of the progress bar
493         int pos = bar_width * progress / 100;
494         // Ensure that progress does not exceed 100%
495         if (progress > 100.0f) {
496             progress = 100.0f;
497         }
498         printf("[%");
499
500         for (int i = 0; i < bar_width; ++i) {
501             if (i < pos) {
502                 printf "=");
503             } else if (i == pos) {
504                 printf ">");
505             } else {
506                 printf " ");
507             }
508         }
509         printf "] %.2f%\r", progress);
510         fflush(stdout); // Flush to update the progress bar immediately
511     }
512
513     printf("\nFile received and written to %s\n", filename);
514
515     // Close the file
516     close(file_fd);
517     return 0;
518 }
519
520 long long getFileSize(char * message) {
521     long long bytes = -1;
522 }
```

```

520 long long getFileSize(char * message) {
521     long long bytes = -1;
522
523     char *start = strrchr(message, '(');
524     if (start != NULL) {
525         char *end = strchr(start, ')');
526         if (end != NULL) {
527             *end = '\0';
528
529             if (sscanf(start + 1, "%lld", &bytes) == 1) {
530                 return bytes;
531             } else {
532                 printf("Falha ao extrair o número de bytes\n");
533             }
534         }
535     }
536     return -1; // Return -1
537 }
538 }
539 }
540 }
```

Appendix II - Part 2 Steps & Commands

Exp 1 - Configure an IP Network

1. Connect the eth1 interface of TUXY3 and TUXY4 to the switch (White Box).
2. Configure the IP addresses for the eth1 interface on each PC:

Unset

```
$ ifconfig eth1 up          # Activate the interface
$ ifconfig eth1 172.16.110.1/24    # For TUXY3
$ ifconfig eth1 172.16.110.254/24   # For TUXY4
```

3. Use ifconfig to check the values (IP and MAC).

4. To verify connectivity between the two computers, use ping (sends packets and waits for a response):

Unset

```
$ ping 172.16.110.254 -c 6  # For TUXY3, the flag limits the number of packets sent
$ ping 172.16.110.1 -c 6    # For TUXY4
```

5. To inspect the Forwarding Table and ARP Table (on TUXY3):

Unset

```
$ route -n    # Forwarding Table
$ arp -a      # ARP Table (ARP translates IP to MAC address, it should display: ?
(172.16.110.1) at <MAC> on ...)
```

6. To delete an entry from the ARP table (on TUXY3):

```
Unset
```

```
$ arp -d 172.16.110.254 # Remove entries associated with TUXY4
```

7. Start Wireshark on TUXY3's eth1 interface.

8. On TUXY3, ping TUXY4:

```
Unset
```

```
$ ping 172.16.110.254 -c 6 # You should observe 6 packets
```

9. You should see packets using the following protocols:

- o ARP -> 2 packets: the first one from the source requesting the MAC address for a given IP, and the second one is the destination's response, sending its MAC address to the source.
- o ICMP -> Used for exchanging information and errors.

Exp 2 - Implement two bridges in a switch

1. Configure eth1 on TUXY2:

```
Unset
```

```
$ ifconfig eth1 up  
$ ifconfig eth1 172.16.111.1/24
```

2. Create two bridges on the switch (TUXY2):

- o First, reset the switch configuration:

```
Unset
```

```
# Open GKTerm and set the baudrate to 115200.  
> admin  
> (password is empty)  
> /system reset-configuration  
> y
```

- o Second, create the two bridges:

Unset

```
> /interface bridge add name=bridge110  
> /interface bridge add name=bridge111
```

- Third, remove the ports where the TUX computers are connected (assuming TUXY2, TUXY3, and TUXY4 are connected to interfaces ether1, ether2, and ether3 of the switch, respectively):

Unset

```
> /interface bridge port remove [find interface=ether2]  
> /interface bridge port remove [find interface=ether3]  
> /interface bridge port remove [find interface=ether4]
```

- Fourth, add the ports (for bridge110, connect TUXY3 and TUXY4; for bridge111, connect only TUXY2):

Unset

```
> /interface bridge port add bridge=bridge110 interface=ether3  
> /interface bridge port add bridge=bridge110 interface=ether4  
> /interface bridge port add bridge=bridge111 interface=ether2
```

- Fifth, to ensure everything is correct, run:

Unset

```
> /interface bridge port print
```

3. Start packet capture on TUXY3.
4. On TUXY3, ping TUXY4 and TUXY2:

Unset

```
$ ping 172.16.110.254 # Ping TUXY4 -> This will work (there is a bridge between  
TUXY3 and TUXY4)  
$ ping 172.16.111.1 # Ping TUXY2 -> "connect: Network is unreachable"
```

5. Stop the captures and start capturing on all TUX computers.
6. On TUXY3, send a broadcast:

Unset

```
$ ping -b 172.16.110.255 # TUXY4 will receive the broadcast, but TUXY2 will not  
(same reason as above)
```

7. Repeat steps 5 and 6, but this time on TUXY2, broadcasting to the subnet ping -b 172.16.111.0/24:

Unset

```
$ ping -b 172.16.111.255 # Neither TUXY4 nor TUXY3 will receive the broadcast  
(same reason as above)
```

Exp 3 - Configure a Router in Linux

1. Activate eth2 on TUXY4:

Unset

```
$ ifconfig eth2 up  
$ ifconfig eth2 172.16.111.253/24
```

2. Add TUXY4 to bridge 111 (connect to the switch as done previously):

Unset

```
> /interface bridge port remove [find interface=ether5]  
> /interface bridge port add bridge=bridge111 interface=ether5
```

3. Enable IP forwarding and disable ICMP echo-ignore-broadcast:

Unset

```
# 1. Enable IP forwarding. This allows the computer to function as a router,  
forwarding packets to other subnets.  
$ echo 1 > /proc/sys/net/ipv4/ip_forward  
# 2. Disable ICMP echo-ignore-broadcast. This allows the PC to respond to  
broadcast messages.  
$ echo 0 > /proc/sys/net/ipv4/icmp_echo_ignore_broadcasts
```

4. Check the IP and MAC addresses for eth1 and eth2 on TUXY4:

Unset

```
$ ifconfig    # The IP address should correspond to each subnet, and the MAC addresses should be different.
```

5. Add routes between TUXY2 and TUXY3 to connect them (the routes pass through TUXY4, which acts as a router):

Unset

```
$ route add -net 172.16.110.0/24 gw 172.16.111.253 # For TUXY2 (<destination subnet> gw <UX's IP>)
$ route add -net 172.16.111.0/24 gw 172.16.110.254 # For TUXY3
# This means: to reach the subnet, the gateway (gw) must be used.
```

6. Check the routes on all three TUX machines:

Unset

```
$ route -n
```

7. Start capturing on TUXY3 and ping the other networks:

Unset

```
$ ping 172.16.110.254 # Must work
$ ping 172.16.111.253 # Must work
$ ping 172.16.111.1 # Must work
```

8. Stop capturing on TUXY3 and start capturing on TUXY4 for both eth1 and eth2.

9. Clear ARP table entries on the TUX machines:

Unset

```
$ arp -d 172.16.111.253 # TUXY2
$ arp -d 172.16.110.254 # TUXY3
$ arp -d 172.16.110.1 # TUXY4
$ arp -d 172.16.111.1 # TUXY4
```

10. On TUXY3, ping TUXY2 and check the logs for both Wireshark instances on TUXY4:

Unset

```
$ ping 172.16.111.1
```

Exp 4 - Configure a Commercial Router and Implement NAT

1. Connect ether1 of the router to the 111.12 strip.
2. Connect ether2 of the router to the switch and add it to bridge111:

```
Unset  
> /interface bridge port remove [find interface=ether6]  
> /interface bridge port add bridge=bridge111 interface=ether6
```

3. Switch the cable from the switch to the router????
4. Connect to the router from TUXY2 using GKTerm:

```
Unset  
# Serial Port: /dev/ttyS0  
# Baudrate: 115200  
# Username: admin  
# Password: <ENTER>
```

5. Reset the configuration:

```
Unset  
> /system reset-configuration
```

6. Configure the router's IP address:

```
Unset  
> /ip address add address=172.16.1.111/24 interface=ether1 # Connect eth1 to the  
RCOM network  
> /ip address add address=172.16.111.254/24 interface=ether2 # Connect to  
bridge111
```

7. Add default routes on the router and TUX machines:

```
Unset  
> /ip route add dst-address=172.16.110.0/24 gateway=172.16.111.253 # Router  
console (Packets for subnet 172.16.110 must go through TUXY4)  
# For some reason, the command below was not used in the screenshots:  
> /ip route add dst-address=0.0.0.0/0 gateway=172.16.1.111 # Router console  
(Associate 0.0.0.0 addresses with the port connecting to the FTP Server)  
## DO NOT USE DEFAULTS = 172.16.1.0/24
```

```
$ route add -net 172.16.1.0/24 gw 172.16.111.254 # TUXY2 (Connect to the router)
$ route add -net 172.16.1.0/24 gw 172.16.110.254 # TUXY3 (Connect to TUXY4,
which in turn has its default route connected to the router)
$ route add -net 172.16.1.0/24 gw 172.16.111.254 # TUXY4 (Connect to the router)
```

8. Verify if TUXY3 can ping other TUX machines and the router (It should work):

Unset

```
$ ping 172.16.110.254 # Ping TUXY4
$ ping 172.16.111.1 # Ping TUXY2
$ ping 172.16.111.254 # Ping Router
```

9. On TUXY2:

- o Disable accept_redirects:

Unset

```
$ echo 0 > /proc/sys/net/ipv4/conf/eth1/accept_redirects
$ echo 0 > /proc/sys/net/ipv4/conf/all/accept_redirects
```

- o Remove the route to the 172.16.110.0/24 network passing through TUXY4:

Unset

```
$ route del -net 172.16.110.0 gw 172.16.111.253 netmask 255.255.255.0
```

- o Ping TUXY3 and observe the packet path using Wireshark (it will go through the router):

Unset

```
$ ping 172.16.110.1 # Uses TUXY2's default gateway (goes to the router) -> Router
forwards to TUXY4 -> TUXY4 forwards to TUXY3
```

- o Perform a traceroute to TUXY3:

Unset

```
$ traceroute -n 172.16.110.1
```

- o Add the route from TUXY2 to TUXY4 back and perform another traceroute:

Unset

```
$ route add -net 172.16.110.0/24 gw 172.16.111.253  
$ traceroute -n 172.16.110.1
```

- Enable accept_redirects:

Unset

```
$ echo 1 > /proc/sys/net/ipv4/conf/eth0/accept_redirects # Shouldn't this be eth1???  
$ echo 1 > /proc/sys/net/ipv4/conf/all/accept_redirects
```

10. On TUXY3, ping the FTP Server (it should work):

Unset

```
$ ping 172.16.1.10
```

11. Disable NAT on the router:

Unset

```
> /ip firewall nat disable 0
```

12. On TUXY4, ping the server (172.16.1.10) and observe what happens (it won't work because packets to 172.16.0.0/16 networks are interpreted as local and routers will drop them):

Unset

```
$ ping 172.16.1.10
```

13. Re-enable NAT:

Unset

```
> /ip firewall nat enable 0
```

Exp 5 - DNS

1. Enable DNS on all TUX machines:

```
Unset
```

```
$ echo "nameserver 10.227.20.3" > /etc/resolv.conf
```

2. Ping a website using its hostname and observe the packets in Wireshark (TUXY2):

```
Unset
```

```
$ ping google.com
```

Exp 6 - TCP connections

1. Start capturing on TUXY3 and run the program.
2. Repeat step 1, but start transferring on TUXY2 at the same time and observe in Wireshark.

Appendix II - Part 2 Wireshark

Exp 1

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|---------------------------|----------------------------|----------------------------|------|
| 1 | 0.000000000 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 2 | 0.431090574 | 192.168.88.1 | 255.255.255.255 | MNDP | 157 5678 → 5678 Len=15 | |
| 3 | 0.431121164 | Routerbo_1c:a3:2a | CDP/VTP/DTP/PAgP/UD... | CDP | 108 Device ID: MikroTik Po | |
| 4 | 2.002153835 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 5 | 4.004308509 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 6 | 6.0006453754 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 7 | 8.008610662 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 8 | 10.010762891 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 9 | 12.012920498 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 10 | 12.076346371 | KYE_10:90:56 | Broadcast | ARP | 42 Who has 172.16.30.254? | |
| 11 | 12.076443939 | Netronix_50:31:bc | KYE_10:90:56 | ARP | 60 172.16.30.254 is at 00: | |
| 12 | 12.076457488 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 13 | 12.076554358 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 14 | 13.107766498 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 15 | 13.107888860 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 16 | 14.015073286 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 17 | 14.131771366 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 18 | 14.131896731 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 19 | 15.155766456 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 20 | 15.155893637 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 21 | 16.017223420 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 22 | 16.179764899 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 23 | 16.179888239 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 24 | 17.179374975 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 25 | 17.179497547 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 26 | 17.237057520 | Netronix_50:31:bc | KYE_10:90:56 | ARP | 60 Who has 172.16.30.1? Te | |
| 27 | 17.237070510 | KYE_10:90:56 | Netronix_50:31:bc | ARP | 42 172.16.30.1 is at 00:c0 | |
| 28 | 18.019396880 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 29 | 18.195760299 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 30 | 18.195877144 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 31 | 19.219766425 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 32 | 19.219894863 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 33 | 20.021535909 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 34 | 20.243765356 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 35 | 20.243881433 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 36 | 21.267771901 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 37 | 21.267907253 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 38 | 22.023688488 | Routerbo_1c:a3:32 | Spanning-tree-(for... STP | 60 RST. Root = 32768/0/c4: | | |
| 39 | 22.268424847 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 40 | 22.268545253 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |
| 41 | 23.283763180 | 172.16.30.1 | 172.16.30.254 | ICMP | 98 Echo (ping) request id | |
| 42 | 23.283890012 | 172.16.30.254 | 172.16.30.1 | ICMP | 98 Echo (ping) reply id | |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0

Exp 2

Apply a display filter ... <Ctrl-/>

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|---------------------------------|---------------------------|------------------------|--------|------|
| 1 | 0.0000000000 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 2 | 2.002201033 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 3 | 4.004409469 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 4 | 6.006603168 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 5 | 8.008809369 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 6 | 10.011010192 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 7 | 12.013212899 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 8 | 14.015415118 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 9 | 16.017610701 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 10 | 18.019812360 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 11 | 20.022006686 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 12 | 22.024206877 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 13 | 24.026411818 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 14 | 26.028605514 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 15 | 28.030810454 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 16 | 30.033011692 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 17 | 32.035214955 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 18 | 33.692539003 | fe80::208:54ff:fe71... ff02::fb | MDNS | 180 Standard query 0x | | |
| 19 | 34.037411652 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 20 | 35.120300725 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 21 | 35.120479520 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 22 | 36.039620711 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 23 | 36.146207633 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 24 | 36.146359400 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 25 | 37.170161481 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 26 | 37.170329312 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 27 | 38.041908273 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 28 | 38.194157514 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 29 | 38.194303833 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 30 | 39.218151032 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 31 | 39.218314742 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 32 | 40.044240253 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 33 | 40.242147204 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 34 | 40.242293942 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 35 | 40.318755166 | Netronix_71:73:ed | ARP | 60 Who has 172.16.60 | | |
| 36 | 40.318770881 | Netronix_50:35:0a | ARP | 42 172.16.60.1 is at | | |
| 37 | 41.266146938 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 38 | 41.266305270 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 39 | 42.046575516 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 40 | 42.290143809 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 41 | 42.290283423 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 42 | 43.314145428 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 43 | 43.314311233 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |
| 44 | 44.048913432 | Routerbo_1c:8d:3b | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 45 | 44.338148654 | 172.16.60.1 | ICMP | 98 Echo (ping) request | | |
| 46 | 44.338306637 | 172.16.60.254 | ICMP | 98 Echo (ping) reply | | |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0

Exp 3

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|------------------------|----------|--------|---------------------|
| 1 | 0.000000000 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 2 | 2.002243753 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 3 | 4.004469975 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 4 | 6.006711771 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 5 | 8.008953357 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 6 | 10.011190334 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 7 | 12.013431570 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 8 | 14.015661980 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 9 | 16.017905800 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 10 | 18.020147244 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 11 | 20.022388408 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 12 | 22.024623566 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 13 | 22.777603183 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 14 | 22.777748524 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 15 | 23.798527766 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 16 | 23.798671222 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 17 | 24.026856279 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 18 | 24.822534614 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 19 | 24.822643149 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 20 | 25.846529170 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 21 | 25.846637076 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 22 | 26.029095347 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 23 | 26.870533434 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 24 | 26.870642178 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 25 | 27.894530504 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 26 | 27.894638340 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 27 | 27.940076679 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 | Who has 172.16.60 |
| 28 | 27.940088412 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | 172.16.60.1 is at |
| 29 | 27.958493304 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | Who has 172.16.60 |
| 30 | 27.958570689 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 | 172.16.60.254 is a |
| 31 | 28.031331970 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 32 | 28.918526945 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 33 | 28.918668585 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 34 | 29.942534351 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 35 | 29.942643584 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 36 | 30.033569920 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 37 | 30.966533097 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 38 | 30.966636882 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 39 | 31.990530306 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 | Echo (ping) request |
| 40 | 31.990638491 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 41 | 32.035799977 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 42 | 34.038034434 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 43 | 36.040276154 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 44 | 38.042520737 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 45 | 40.044753516 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 46 | 42.046986784 | Routerbo_1c:8d:3b | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|---------------|-------------------|------------------------|----------|--------|---------------------|
| 47 | 44.049220260 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 48 | 44.217734794 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 49 | 44.217881462 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 50 | 45.238537081 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 51 | 45.238643241 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 52 | 46.051465121 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 53 | 46.262535337 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 54 | 46.262635350 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 55 | 46.343698395 | 0.0.0.0 | 255.255.255.255 | MNDP | 159 | 5678 → 5678 Len=1: |
| 56 | 46.343729894 | Routerbo_1c:8d:3b | CDP/VTP/DTP/PAgP/UD... | CDP | 93 | Device ID: MikroT |
| 57 | 46.343778923 | Routerbo_1c:8d:3b | LLDP_Multicast | LLDP | 110 | MA/c4:ad:34:1c:8d |
| 58 | 47.286534290 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 59 | 47.286639263 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 60 | 48.053710680 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 61 | 48.310531288 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 62 | 48.310637937 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 63 | 49.334531918 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 64 | 49.334665944 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 65 | 50.055499681 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 66 | 50.358533665 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 67 | 50.358642409 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 68 | 51.382531431 | 172.16.60.1 | 172.16.61.253 | ICMP | 98 | Echo (ping) request |
| 69 | 51.382638079 | 172.16.61.253 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 70 | 52.057727639 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 71 | 54.059963698 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 72 | 54.564092719 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 | Who has 172.16.60 |
| 73 | 54.564107665 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | 172.16.60.1 is at |
| 74 | 56.062207998 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 75 | 58.064453694 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 76 | 58.474034376 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 77 | 58.474426958 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 78 | 59.482524852 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 79 | 59.482771813 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 80 | 60.0666687307 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 81 | 60.502530311 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 82 | 60.502747519 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 83 | 61.526529752 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 84 | 61.526765329 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 85 | 62.068922247 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 86 | 62.550530799 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 87 | 62.550743538 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 88 | 63.574530100 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 89 | 63.574770077 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 90 | 64.071159422 | Routerbo_1c:8d:3b | Spanning-tree-(for... | STP | 60 | RST. Root = 32768, |
| 91 | 64.602526177 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 92 | 64.602739125 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0

Exp 4

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|---------------|-------------------|---------------------------|-----------------------|------------------------|------|
| 1 | 0.000000000 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 2 | 2.002136126 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 3 | 4.004234746 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 4 | 6.006361093 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 5 | 8.008508114 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 6 | 10.010603801 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 7 | 12.012742371 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 8 | 14.014870255 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 9 | 16.016959167 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 10 | 18.0009091971 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 11 | 20.011222230 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 12 | 21.176807045 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 Echo (ping) request | |
| 13 | 21.176938417 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 14 | 22.013305415 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 15 | 22.197551818 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 Echo (ping) request | |
| 16 | 22.197662518 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 17 | 23.221558560 | 172.16.60.1 | 172.16.60.254 | ICMP | 98 Echo (ping) request | |
| 18 | 23.221683158 | 172.16.60.254 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 19 | 24.015431274 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 20 | 26.017552803 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 21 | 26.325518345 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 Who has 172.16.60 | |
| 22 | 26.325618079 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 172.16.60.254 is at | |
| 23 | 26.429105194 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 Who has 172.16.60 | |
| 24 | 26.429123911 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 172.16.60.1 is at | |
| 25 | 28.019664135 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 26 | 30.021762896 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 27 | 32.023901047 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 28 | 34.026002741 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 29 | 36.028150321 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 30 | 38.030299438 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 31 | 38.698640475 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 Echo (ping) request | |
| 32 | 38.698919564 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 33 | 39.342708574 | 0.0.0.0 | 255.255.255.255 | MNDP | 159 5678 → 5678 Len=1 | |
| 34 | 39.342738047 | Routerbo_1c:8d:2d | CDP/VTP/DTP/PAgP/UD... | CDP | 93 Device ID: MikroT | |
| 35 | 39.342785609 | Routerbo_1c:8d:2d | LLDP_Multicast | LLDP | 110 MA/c4:ad:34:1c:8d | |
| 36 | 39.701553413 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 Echo (ping) request | |
| 37 | 39.701794927 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 38 | 40.032401132 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 39 | 40.725557571 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 Echo (ping) request | |
| 40 | 40.725791332 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |
| 41 | 42.034105914 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 42 | 44.036672197 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 43 | 46.038756501 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 44 | 48.040900938 | Routerbo_1c:8d:2d | Spanning-tree-(for... STP | 60 RST. Root = 32768, | | |
| 45 | 48.921109471 | 172.16.60.1 | 172.16.61.254 | ICMP | 98 Echo (ping) request | |
| 46 | 48.921451068 | 172.16.61.254 | 172.16.60.1 | ICMP | 98 Echo (ping) reply | |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0

File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

Apply a display filter ... <Ctrl-/>

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|------------------------|----------|--------|-------------------------|
| 24 | 26.429123911 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | 172.16.60.1 is at |
| 25 | 28.019664135 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 26 | 30.021762896 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 27 | 32.023901047 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 28 | 34.026002741 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 29 | 36.028150321 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 30 | 38.030299438 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 31 | 38.698640475 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 32 | 38.698919564 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 33 | 39.342708574 | 0.0.0.0 | 255.255.255.255 | MNDP | 159 | 5678 → 5678 Len=13 |
| 34 | 39.342738047 | Routerbo_1c:8d:2d | CDP/FTP/DTP/PAgP/UD... | CDP | 93 | Device ID: MikroT |
| 35 | 39.342785609 | Routerbo_1c:8d:2d | LLDP_Multicast | LLDP | 110 | MA/c4:ad:34:1c:8d |
| 36 | 39.701553413 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 37 | 39.701794927 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 38 | 40.032401132 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 39 | 40.725557571 | 172.16.60.1 | 172.16.61.1 | ICMP | 98 | Echo (ping) request |
| 40 | 40.725791332 | 172.16.61.1 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 41 | 42.034105914 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 42 | 44.036672197 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 43 | 46.038756501 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 44 | 48.040900938 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 45 | 48.921109471 | 172.16.60.1 | 172.16.61.254 | ICMP | 98 | Echo (ping) request |
| 46 | 48.921451068 | 172.16.61.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 47 | 49.941550014 | 172.16.60.1 | 172.16.61.254 | ICMP | 98 | Echo (ping) request |
| 48 | 49.941820372 | 172.16.61.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 49 | 50.043032595 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 50 | 50.965549143 | 172.16.60.1 | 172.16.61.254 | ICMP | 98 | Echo (ping) request |
| 51 | 50.965819082 | 172.16.61.254 | 172.16.60.1 | ICMP | 98 | Echo (ping) reply |
| 52 | 52.045142461 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 53 | 54.047266644 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 54 | 54.077108711 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 42 | Who has 172.16.60 |
| 55 | 54.077124285 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | 172.16.60.1 is at |
| 56 | 56.039405316 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 57 | 57.129764420 | 172.16.60.1 | 193.136.152.72 | NTP | 90 | NTP Version 4, cl |
| 58 | 57.130067604 | 172.16.61.254 | 172.16.60.1 | ICMP | 118 | Destination unreachable |
| 59 | 58.041454349 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 60 | 60.043547942 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 61 | 62.045673243 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 62 | 62.165517761 | Netronix_50:35:0a | Netronix_71:73:ed | ARP | 42 | Who has 172.16.60 |
| 63 | 62.165615749 | Netronix_71:73:ed | Netronix_50:35:0a | ARP | 60 | 172.16.60.254 is a |
| 64 | 64.047732614 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 65 | 66.049787026 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 66 | 68.051970715 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 67 | 70.053931050 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |
| 68 | 72.056238080 | Routerbo_1c:8d:2d | Spanning-tree-(for...) | STP | 60 | RST. Root = 32768, |

Frame 1: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface eth1, id 0